

Naive Bayes Project

Matthew Lueder

February 9, 2017

Load Packages

```
require('Rcpp')

## Loading required package: Rcpp
## Warning: package 'Rcpp' was built under R version 3.2.4

require('inline')

## Loading required package: inline
## Warning: package 'inline' was built under R version 3.2.5

##
## Attaching package: 'inline'

## The following object is masked from 'package:Rcpp':
##
##   registerPlugin

require('microbenchmark')

## Loading required package: microbenchmark
## Warning: package 'microbenchmark' was built under R version 3.2.5

require('caret')

## Loading required package: caret
## Warning: package 'caret' was built under R version 3.2.4

## Loading required package: lattice
## Loading required package: ggplot2
## Warning: package 'ggplot2' was built under R version 3.2.4
```

Read in training data

Training data will be loaded into a list, with each element being a list of words for that training example. The first word for each line of input data is the label. This is removed and stored in the `lbls` vector. A vector containing each class is created (`classes`).

```
# Set working directory to project folder
setwd("C:/Users/Matthew/Dropbox/Machine Learning/Projects/Project2")

training = scan("T_S.data", what = character(), fill = T, sep = '\n')
training = lapply(training, strsplit, split = "[ ]")
lbls = unlist(lapply(training, function(l) unlist(l)[1]))
classes = unique(lbls)
training = lapply(training, function(l) unlist(l)[-1]) # Remove first
element (the label)
```

Build Vocabulary

The vocabulary is the set of words which have appeared at least once in the training data

```
vocab = unique(unlist(training))
```

Calculate P(Class) for each class

We create a vector of the probabilities of seeing each class based on the frequency of appearance of that class in the training data.

$$P(Class_i) = \frac{\text{Num. of training examples of in class } i}{\text{Total number of training examples}}$$

```
P.Class = sapply(classes, function(class) length(lbls[lbls==class])) /
length(lbls)
```

Combine training examples of same class

Create a list with one index for each class and assign each index all words contained in all training examples of the respective class

```
Texts = sapply(classes, function(class) unlist(training[lbls == class]))
```

Count words for each class

- 1) Create an empty matrix for counting the words of each class
- 2) Create a C++ function to words quickly
- 3) Apply the C++ function A C++ function was used because counting words with an R loop took an extraordinary amount of time. We compare the performance of this C++ function to R in the next code chunk.

```
# 1)
counts = vector(mode = "integer", length = length(vocab))
names(counts) = vocab
counts = t(sapply(classes, function(class) counts))
```

```
# 2)
funcSrc <- '
#include <map>
```

```

IntegerMatrix counts(counts_in);
CharacterVector classes(classes_in);
CharacterVector vocab(vocab_in);
List texts(texts_in);

// Create map to look up position of words/class indices
std::map<std::string, int> wordIndex;
std::map<std::string, int> classIndex;
for (int i=0; i != vocab.size(); ++i)
{
    wordIndex[as<std::string>(vocab[i])] = i;
}
for (int i=0; i != classes.size(); ++i)
{
    classIndex[as<std::string>(classes[i])] = i;
}

// Count words
for (int i=0; i != classes.size(); ++i)
{
    CharacterVector text =
Rcpp::as<CharacterVector>(texts[as<std::string>(classes[i])]);

    for (int j=0; j != text.size(); ++j)
    {
        counts( classIndex[as<std::string>(classes[i])],
wordIndex[as<std::string>(text[j])] )++;
    }
}
return counts;
'

count.words <- cxxfunction(sig = signature(counts_in="integer",
classes_in="character", texts_in="list", vocab_in="character"), funcSrc,
plugin = "Rcpp")

# 3)
counts = count.words(counts, classes, Texts, vocab)

rm(funcSrc)

```

Compare performance of C++ and R

Shows why C++ was necessary, and how much it improved performance. First I reduce the dataset to 1/1000th of the size to allow R function to complete in reasonable time

```

Texts.short = lapply(Texts, function(words)
words[0:as.integer(length(words)/1000)])

count.words.inR <- function(counts, classes, Texts) {
    for(class in classes)

```

```

    {
      for(word in unlist(Texts[class]))
      {
        counts[class, word] = counts[class, word] + 1
      }
    }
  }

microbenchmark( count.words(counts, classes, Texts.short, vocab),
                 count.words.inR(counts, classes, Texts.short),
                 times = 10 )

## Unit: milliseconds
##                               expr          min          lq
## count.words(counts, classes, Texts.short, vocab)  57.52659   58.92327
## count.words.inR(counts, classes, Texts.short) 7770.37234 8155.74572
##      mean      median        uq      max neval cld
##  59.93083   60.13951   60.46035   63.7283    10  a
## 8495.30335 8359.92956 8499.68593 10362.9070    10  b

rm(count.words.inR, Texts.short)

```

Calculate word probabilities for each class

$$P(\text{Word}_k | \text{Class}_i) = (n_k + 1) / (n + |\text{Vocabulary}|)$$

n = total number of word positions in class i n_k = number of times Word k occurs in class

```

n = apply(counts, 1, sum)
n.v = n + length(vocab) # = n + |Vocabulary|
P.words = (counts + 1) / n.v
Log.P.words = log(P.words)

rm(n, n.v, P.words)

```

Read in validation data

```

validation = scan("V_S.data", what = character(), fill = T, sep = '\n')
validation = lapply(validation, strsplit, split = "[ ]")
v.lbls = unlist(lapply(validation, function(l) unlist(l)[1]))
validation = lapply(validation, function(l) unlist(l)[-1]) # Remove first
element (the label)

# Remove words not in vocabulary
validation = lapply(validation, function(input) input[unlist(input) %in%
vocab])

```

Perform classification

```

classify <- function(example) {
  which.max(rowSums(cbind(Log.P.words[, unlist(example)], log(P.Class))))
}

```

```
PredictedClasses = factor(names(sapply(validation, classify)))
```

Collect metrics / Visualize results

```
confusionMatrix(PredictedClasses, factor(v.lbls))
```

```
## Confusion Matrix and Statistics
```

```
##
```

```
##           Reference
```

```
## Prediction      atheism autos baseball christianity cryptology
```

```
## atheism          242      0         0           8           0
```

```
## autos             0     369         0           0           1
```

```
## baseball          0      2        364           0           0
```

```
## christianity       40      0         2        376           0
```

```
## cryptology        1      3         2           0        376
```

```
## electronics       0      4         0           0           4
```

```
## forsale           0      2         2           0           0
```

```
## graphics          0      1         0           1           2
```

```
## guns              4      2         1           2           5
```

```
## hockey            1      0        17           0           0
```

```
## mac                0      0         0           0           0
```

```
## medicine          3      0         0           1           1
```

```
## mideastpolitics   8      2         2           2           0
```

```
## motorcycles       1      3         1           0           0
```

```
## mswindows         0      0         0           1           2
```

```
## pc                 0      0         0           0           2
```

```
## politics          6      5         3           3           2
```

```
## religion          10      0         0           3           0
```

```
## space              3      2         1           1           0
```

```
## xwindows          0      0         2           0           1
```

```
##           Reference
```

```
## Prediction      electronics forsale graphics guns hockey mac medicine
```

```
## atheism          0         0         0         0         0         0         7
```

```
## autos            9        19         0         0         0         3         3
```

```
## baseball         0         2         0         0         1         1         0
```

```
## christianity       2         2         1         1         3         0        10
```

```
## cryptology       48         2        14         5         1         6         0
```

```
## electronics      266        10         4         0         0        12         4
```

```
## forsale           1       261         0         1         0         5         0
```

```
## graphics         12         5       316         0         0         9         8
```

```
## guns              0         6         0      335         0         2         5
```

```
## hockey            0         1         0         0       389         0         0
```

```
## mac               6        22        10         0         0      293         1
```

```
## medicine          6         5         2         0         0         4      336
```

```
## mideastpolitics   3         0         1         1         1         0         5
```

```
## motorcycles       5         5         0         1         0         1         0
```

```
## mswindows         4         2         5         0         0         7         1
```

```
## pc                28        38        10         0         0        34         1
```

```
## politics          0         3         0        15         2         2        12
```

```
## religion          0         0         1         3         0         0         0
```

```

##      space      3      5      8      2      2      4      3
##      xwindows    0      2     17     0      0      2      0
##
##      Reference
## Prediction  mideastpolitics motorcycles mswindows  pc  politics
##      atheism      7      0      2      0      6
##      autos      1     14      0      0      0
##      baseball    1      2      1      0      0
##      christianity  4      1      2      0      2
##      cryptology   4      0     19      6      5
##      electronics  0      1      1     26      0
##      forsale      0      1      0      5      0
##      graphics     0      0     34      8      0
##      guns         6      1      0      0     100
##      hockey       0      0      0      0      0
##      mac          0      0      8     30      0
##      medicine     0      1      1      0      1
##      mideastpolitics 338      0      0      0      3
##      motorcycles  2     372      0      2      0
##      mswindows    0      0     245     23      0
##      pc           0      0      33    290      0
##      politics     13      4     13      0     184
##      religion     0      0      1      0      0
##      space        0      1      3      1      9
##      xwindows     0      0     30      1      0
##
##      Reference
## Prediction  religion space xwindows
##      atheism    43      1      0
##      autos      0      0      0
##      baseball    0      0      1
##      christianity 68      1      0
##      cryptology   0      1      9
##      electronics  0      5      1
##      forsale      0      0      1
##      graphics     4      8     42
##      guns        21      1      2
##      hockey       0      0      0
##      mac          0      0      6
##      medicine     2      4      1
##      mideastpolitics 3      2      2
##      motorcycles  0      0      1
##      mswindows    0      0      8
##      pc           0      0      4
##      politics     6     15      2
##      religion     98      1      0
##      space        6    352      4
##      xwindows     0      3    308
##
## Overall Statistics
##
## Accuracy : 0.8117

```

```

##          95% CI : (0.8027, 0.8205)
##      No Information Rate : 0.053
##      P-Value [Acc > NIR] : < 2.2e-16
##
##          Kappa : 0.8017
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##          Class: atheism Class: autos Class: baseball
## Sensitivity          0.75862      0.93418      0.91688
## Specificity          0.98973      0.99299      0.99846
## Pos Pred Value       0.76582      0.88067      0.97067
## Neg Pred Value       0.98932      0.99634      0.99539
## Prevalence           0.04238      0.05248      0.05274
## Detection Rate       0.03215      0.04902      0.04836
## Detection Prevalence 0.04198      0.05567      0.04982
## Balanced Accuracy     0.87418      0.96358      0.95767
##
##          Class: christianity Class: cryptology
## Sensitivity          0.94472      0.94949
## Specificity          0.98050      0.98233
## Pos Pred Value       0.73010      0.74900
## Neg Pred Value       0.99686      0.99715
## Prevalence           0.05288      0.05261
## Detection Rate       0.04995      0.04995
## Detection Prevalence 0.06842      0.06669
## Balanced Accuracy     0.96261      0.96591
##
##          Class: electronics Class: forsale Class: graphics
## Sensitivity          0.67684      0.66923      0.81234
## Specificity          0.98991      0.99748      0.98123
## Pos Pred Value       0.78698      0.93548      0.70222
## Neg Pred Value       0.98233      0.98220      0.98968
## Prevalence           0.05221      0.05181      0.05168
## Detection Rate       0.03534      0.03468      0.04198
## Detection Prevalence 0.04491      0.03707      0.05978
## Balanced Accuracy     0.83338      0.83335      0.89678
##
##          Class: guns Class: hockey Class: mac Class: medicine
## Sensitivity          0.92033      0.97494      0.76104      0.84848
## Specificity          0.97794      0.99733      0.98838      0.99551
## Pos Pred Value       0.67951      0.95343      0.77926      0.91304
## Neg Pred Value       0.99588      0.99860      0.98713      0.99162
## Prevalence           0.04836      0.05301      0.05115      0.05261
## Detection Rate       0.04451      0.05168      0.03893      0.04464
## Detection Prevalence 0.06550      0.05420      0.04995      0.04889
## Balanced Accuracy     0.94914      0.98614      0.87471      0.92200
##
##          Class: mideastpolitics Class: motorcycles
## Sensitivity          0.89894      0.93467
## Specificity          0.99511      0.99691
## Pos Pred Value       0.90617      0.94416
## Neg Pred Value       0.99469      0.99635

```

## Prevalence	0.04995	0.05288
## Detection Rate	0.04491	0.04942
## Detection Prevalence	0.04955	0.05234
## Balanced Accuracy	0.94702	0.96579
##	Class: mswindows	Class: pc
## Sensitivity	0.62341	0.73980
## Specificity	0.99257	0.97898
## Pos Pred Value	0.82215	0.65909
## Neg Pred Value	0.97953	0.98561
## Prevalence	0.05221	0.05208
## Detection Rate	0.03255	0.03853
## Detection Prevalence	0.03959	0.05846
## Balanced Accuracy	0.80799	0.85939
##	Class: religion	Class: space
## Sensitivity	0.39044	0.89340
## Specificity	0.99739	0.99187
## Pos Pred Value	0.83761	0.85854
## Neg Pred Value	0.97935	0.99410
## Prevalence	0.03335	0.05234
## Detection Rate	0.01302	0.04676
## Detection Prevalence	0.01554	0.05447
## Balanced Accuracy	0.69391	0.94263

Overall accuracy = 81.04% Let's see if we can do better.

Eliminate words which have a probability with a standard deviation lower than a certain amount

```
Log.P.words.store = Log.P.words
accuracy = c()

for (sdp in seq(.1,.9,.1))
{
  Log.P.words = Log.P.words.store
  sdForEachWord = apply(exp(Log.P.words), 2, sd)
  wordsToElim = which(sdForEachWord < quantile(sdForEachWord, sdp))
  Log.P.words = Log.P.words[, -wordsToElim]
  vs.vocab = vocab[-wordsToElim]
  vs.validation = lapply(validation, function(input) input[unlist(input) %in%
vs.vocab])
  PredictedClasses = factor(names(sapply(vs.validation, classify)))
  accuracy = c( accuracy, confusionMatrix(PredictedClasses,
factor(v.lbls))$overall["Accuracy"])
}


rm(sdp, sdForEachWord, wordsToElim, vs.validation, vs.vocab)
```

See if we did any better

accuracy


```
## Accuracy Accuracy Accuracy Accuracy Accuracy Accuracy Accuracy
## 0.8113458 0.8108144 0.8102830 0.8100173 0.8078916 0.8057659 0.8028431
## Accuracy Accuracy
## 0.7959346 0.7737478
```

```
plot(seq(10,90,10), accuracy, xlab = "Percent of Words Removed Based on
Standard Deviation")
```



Percent of Words Removed Based on Standard Deviation

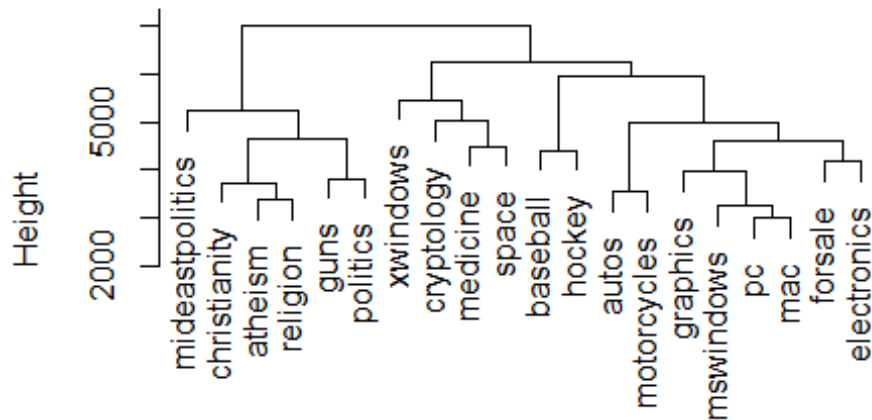
Removing features based on standard deviation hurts the performance of the classifier after removing more than 20% of the words. Removing 10% helped the performance of the classifier but only an extremely small, insignificant amount.

The classes seemed to be related, ex. Christianity, atheism, religion. Maybe it makes sense to group together certain classes. Let's see what hierarchical clustering tells us.

Perform hierarchical clustering

```
clusters = hclust(dist(Log.P.words, method = "manhattan"))
plot(clusters)
```

Cluster Dendrogram



```
dist(Log.P.words, method = "manhattan")
hclust (*, "complete")
```

```
Log.P.words = Log.P.words.store
```

Looking at the results we see that a lot of our clusters make sense. We can use these results to guide us in how we group classes. Christianity, atheism, and religion are closely clustered so we will group them into a single class, religion. xwindows, graphics, mswindows, pc, and mac are closely clustered, so we will group them into a computing class. Cryptology, medicine and space are clustered together, so we will group them into a class named science. Baseball and hockey are clustered together, we will group them in a class named sports. Autos and motorcycles are clustered together, so we will group them in a class called vehicles.

Restructure training data

```
reclassify <- function(class){
  if (class %in% c("atheism", "christianity")) {
    return("religion")
  }
  if (class %in% c("xwindows", "graphics", "mswindows", "pc", "mac")) {
    return("computing")
  }
  if (class %in% c("cryptology", "medicine", "space")) {
    return("science")
  }
  if (class %in% c("baseball", "hockey")) {
    return("sports")
  }
  if (class %in% c("autos", "motorcycles")) {
```

```

    return("vehicles")
}
return(class)
}

lbls = sapply(lbls, reclassify)
classes = unique(lbls)

P.Class = sapply(classes, function(class) length(lbls[lbls==class])) /
length(lbls)
Texts = sapply(classes, function(class) unlist(training[lbls == class]))

counts = vector(mode = "integer", length = length(vocab))
names(counts) = vocab
counts = t(sapply(classes, function(class) counts))
counts = count.words(counts, classes, Texts, vocab)

n = apply(counts, 1, sum)
n.v = n + length(vocab) # = n + |Vocabulary|
P.words = (counts + 1) / n.v
Log.P.words = log(P.words)
rm(n, n.v, P.words)

```

Restructure validation data

```
v.lbls = sapply(v.lbls, reclassify)
```

Perform classification with restructured data

And print results.

```

PredictedClasses = factor(names(sapply(validation, classify)))
confusionMatrix(PredictedClasses, factor(v.lbls))

```

```
## Confusion Matrix and Statistics
```

```
##
##              Reference
## Prediction    computing electronics forsale guns mideastpolitics
## computing      1846          138      132    1          1
## electronics     14          175       5    0          0
## forsale         6           1      188    0          0
## guns            1           0       1  323          3
## mideastpolitics 0           0       0    1        312
## politics        3           0       0   10         10
## religion        9           3       3   12         40
## science        63          48      14   16          5
## sports          1           0       3    0          2
## vehicles        8          28      44    1          3
##
##              Reference
## Prediction    politics religion science sports vehicles
## computing      0           9      28    5          4

```

```

##      electronics      0      0      6      0      1
##      forsale          0      0      0      0      1
##      guns            94     16      5      0      0
##      mideastpolitics   2      7      3      1      1
##      politics        166     4      4      1      1
##      religion         26     906     44      5      4
##      science          22     21    1086      8      5
##      sports           0      1      1     772      2
##      vehicles         0      4      9      4     774
##
## Overall Statistics
##
##              Accuracy : 0.8699
##              95% CI : (0.8621, 0.8775)
##      No Information Rate : 0.2592
##      P-Value [Acc > NIR] : < 2.2e-16
##
##              Kappa : 0.8467
##      McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##              Class: computing Class: electronics Class: forsale
## Sensitivity          0.9462          0.44529      0.48205
## Specificity          0.9430          0.99636      0.99888
## Pos Pred Value       0.8530          0.87065      0.95918
## Neg Pred Value       0.9804          0.97024      0.97245
## Prevalence           0.2592          0.05221      0.05181
## Detection Rate       0.2453          0.02325      0.02498
## Detection Prevalence 0.2875          0.02670      0.02604
## Balanced Accuracy    0.9446          0.72082      0.74047
##
##              Class: guns Class: mideastpolitics Class: politics
## Sensitivity          0.88736          0.82979      0.53548
## Specificity          0.98325          0.99790      0.99543
## Pos Pred Value       0.72912          0.95413      0.83417
## Neg Pred Value       0.99421          0.99111      0.98035
## Prevalence           0.04836          0.04995      0.04119
## Detection Rate       0.04291          0.04145      0.02205
## Detection Prevalence 0.05885          0.04344      0.02644
## Balanced Accuracy    0.93530          0.91384      0.76546
##
##              Class: religion Class: science Class: sports
## Sensitivity          0.9360          0.9157      0.9698
## Specificity          0.9777          0.9681      0.9985
## Pos Pred Value       0.8612          0.8432      0.9872
## Neg Pred Value       0.9904          0.9840      0.9964
## Prevalence           0.1286          0.1576      0.1058
## Detection Rate       0.1204          0.1443      0.1026
## Detection Prevalence 0.1398          0.1711      0.1039
## Balanced Accuracy    0.9568          0.9419      0.9842
##
##              Class: vehicles

```

```
## Sensitivity          0.9760
## Specificity          0.9850
## Pos Pred Value      0.8846
## Neg Pred Value      0.9971
## Prevalence          0.1054
## Detection Rate      0.1028
## Detection Prevalence 0.1162
## Balanced Accuracy    0.9805
```

Overall accuracy = 0.8699, an improvement. Forsale and electronics were difficult to classify.

It is suggested in literature to remove class prior probabilities when classifying text documents. Let's see what happens when we do this.

Perform classification again, but drop P(Class) from the final equation

```
zeros = rep(0, length(classes))
classify2 <- function(example) {
  which.max(rowSums(cbind(Log.P.words[,unlist(example)], zeros)))
}
```

```
PredictedClasses = factor(names(sapply(validation, classify2)))
confusionMatrix(PredictedClasses, factor(v.lbls))
```

```
## Confusion Matrix and Statistics
```

```
##
##              Reference
## Prediction      computing electronics forsale guns mideastpolitics
## computing      1841          126      121    1              0
## electronics     15          190        8    0              0
## forsale         7           1       197    0              0
## guns            2           0        2  324              3
## mideastpolitics 0           0        0   1             314
## politics        4           0        0  10             13
## religion        9           3        3  11             37
## science        63          46       12  16             4
## sports         2           0        3   0              2
## vehicles       8           27       44   1              3
```

```
##              Reference
## Prediction      politics religion science sports vehicles
## computing       0           8       28    5           4
## electronics     0           0        6    0           1
## forsale         0           0        0    0           2
## guns            93          17        8    0           0
## mideastpolitics 3           7        5    1           1
## politics       170          4        5    1           2
## religion       24          906       42    6           4
## science       20           21     1082    7           4
## sports         0           1        1   772          2
## vehicles       0           4        9    4          773
```

```

##
## Overall Statistics
##
##           Accuracy : 0.8727
##           95% CI : (0.865, 0.8802)
##           No Information Rate : 0.2592
##           P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.8502
##           Mcnemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: computing Class: electronics Class: forsale
## Sensitivity           0.9436           0.48346           0.50513
## Specificity           0.9475           0.99579           0.99860
## Pos Pred Value        0.8627           0.86364           0.95169
## Neg Pred Value        0.9796           0.97222           0.97363
## Prevalence            0.2592           0.05221           0.05181
## Detection Rate        0.2446           0.02524           0.02617
## Detection Prevalence  0.2835           0.02923           0.02750
## Balanced Accuracy     0.9455           0.73963           0.75186
##
##           Class: guns Class: mideastpolitics Class: politics
## Sensitivity           0.89011           0.83511           0.54839
## Specificity           0.98255           0.99748           0.99460
## Pos Pred Value        0.72160           0.94578           0.81340
## Neg Pred Value        0.99435           0.99138           0.98087
## Prevalence            0.04836           0.04995           0.04119
## Detection Rate        0.04305           0.04172           0.02259
## Detection Prevalence  0.05965           0.04411           0.02777
## Balanced Accuracy     0.93633           0.91629           0.77149
##
##           Class: religion Class: science Class: sports
## Sensitivity           0.9360           0.9123           0.9698
## Specificity           0.9788           0.9696           0.9984
## Pos Pred Value        0.8670           0.8486           0.9860
## Neg Pred Value        0.9904           0.9834           0.9964
## Prevalence            0.1286           0.1576           0.1058
## Detection Rate        0.1204           0.1437           0.1026
## Detection Prevalence  0.1388           0.1694           0.1040
## Balanced Accuracy     0.9574           0.9409           0.9841
##
##           Class: vehicles
## Sensitivity           0.9748
## Specificity           0.9851
## Pos Pred Value        0.8855
## Neg Pred Value        0.9970
## Prevalence            0.1054
## Detection Rate        0.1027
## Detection Prevalence  0.1160
## Balanced Accuracy     0.9800

```

Overall accuracy = 0.8727, a slight improvement.